

May 26, 2025

Graph Structure Learning

Alessandro Manenti

Graph Machine Learning Group (gmlg.ch)

The Swiss AI Lab IDSIA

Università della Svizzera italiana






Introduction

Introduction

GNNs use an adjacency matrix A as an effective **inductive bias**.

:(A might be **unknown** or of **coarsely available**

Some examples:



Can we learn relationships from data?

Introduction

- 😊 It is possible to learn relations from data

Graph Structure Learning (GSL) investigates methods to infer relational structures from data.

GSL effectiveness depends on:

1. The presence of a "true" underlying relational structure.
2. The number of available data


The Transformer learns relational structures from data too:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad \text{with: } Q/K/V = W_Q/W_K/W_V \cdot \mathbf{X}$$

Q: Where is the relational structure here?

Attention mechanism

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad \text{with: } Q/K/V = W_Q/W_K/W_V \cdot \mathbf{X}$$



Overview


Using original structure or Adjacency matrix initialization	Graph structure learning	Transformer-based techniques
<ul style="list-style-type: none">Pre-processing techniques used to infer an initial, static topology	<ul style="list-style-type: none">Techniques that parametrize and optimize the structure to solve a task	<ul style="list-style-type: none">Techniques based on the attention mechanism
Limited data		→ Abundant data
Computationally efficient		→ Computationally expensive

- For further reading, refer to [1], [2]

[1] Zhiyao et al., “Opengsl: A comprehensive benchmark for graph structure learning” 2024.

[2] Fatemi et al., “Ugsl: A unified framework for benchmarking graph structure learning” 2023.

General GSL Framework



- Input: $\mathbf{X} \in \mathbb{R}^{N \times D}$ and, optionally, an initial adjacency matrix $\mathbf{A}^{(0)} \in \mathbb{R}^{N \times N}$
- Trainable modules: Edge Scorer and GNN
- Loss function: Usually designed to solve a (self-)supervised task

Structure initialization techniques

Structure initialization techniques

- Extract (or modify) an adjacency matrix **independently** from the downstream task.
 - Different techniques rely on different **assumptions**.
- 😊 Topological structures obtained from this pre-processing can be used as **initialization** for the GSL edge scorer.

Some examples include:

1. **Pearson** Correlation.
2. **Granger** causality.
3. Pairwise **input similarity**.
4. **Dirichlet Energy Minimization**.
5. **Rewiring** techniques (if initial $A^{(0)}$ given).

Pearson correlation

The Pearson correlation coefficient is a measure of the **linear relationship** between two variables.

$$\rho \equiv \frac{\text{Cov}(\mathbf{X}_i, \mathbf{X}_j)}{\sigma_{\mathbf{X}_i} \sigma_{\mathbf{X}_j}}$$

For real-world data the formula is:

$$\hat{\rho} = \frac{\sum_{d=1}^D (\mathbf{X}_{i,d} - \bar{\mathbf{X}}_i)(\mathbf{X}_{j,d} - \bar{\mathbf{X}}_j)}{\sqrt{\sum_{d=1}^D (\mathbf{X}_{i,d} - \bar{\mathbf{X}}_i)^2 \sum_{d=1}^D (\mathbf{X}_{j,d} - \bar{\mathbf{X}}_j)^2}}$$

An adjacency matrix \mathbf{A} can be built from $\hat{\rho}$.

Pearson correlation

- ρ is a **normalized** value: $-1 \leq \rho \leq 1$
- The **magnitude** of ρ indicates the strength of the relationship,
- The **sign** indicates its direction.
- Be aware that it is not perfect! (see Figure)




Figure 1: Pearson correlation for different sets of (x, y) points. Image from Wikipedia

Granger causality

For Granger causality, we restrict \mathbf{X} to be a set of time series.

- Granger causality test exists if time series \mathbf{X}_i "causes" time series \mathbf{X}_j .
- Test whether past values of \mathbf{X}_i contain useful information for predicting \mathbf{X}_j , beyond the information contained in past values of \mathbf{X}_j alone.

Build two linear models:

Restricted model (without \mathbf{X}_j)

$$\mathbf{X}_{i,t} = \alpha_0 + \sum_{a=1}^p \alpha_a \mathbf{X}_{i,t-a} + \epsilon_t$$

Unrestricted model (with \mathbf{X}_j)

$$\mathbf{X}_{i,t} = \alpha_0 + \sum_{a=1}^p \alpha_a \mathbf{X}_{i,t-a} + \sum_{b=1}^p \gamma_b \mathbf{X}_{j,t-b} + \eta_t$$

The Granger causality test assesses whether \mathbf{X}_j helps to predict \mathbf{X}_i .

Granger causality

Formulate the **null** hypothesis H_0 and **alternative** hypothesis H_1 :

$$H_0 : \gamma_1 = \gamma_2 = \cdots = \gamma_p = 0$$

$$H_1 : \text{At least one } \gamma_b \neq 0 \text{ for some } b \in \{1, 2, \dots, p\}$$

H_0 : none of the past values of \mathbf{X}_j contain linear predictive information about the current value of \mathbf{X}_i .

To test H_0 , compare the fit of the restricted and unrestricted models. This is typically done using an **F-test**:

1. Compute the residual sum of squares (RSS) for both the **restricted** model (RSS_R) and the **unrestricted** model (RSS_U)
2. Compute the F-statistic:

$$\frac{(RSS_R - RSS_U)/p}{RSS_U/(T - 2p - 1)}$$

Under H_0 , the F-statistic follows an F-distribution with p and $(T - 2p - 1)$ degrees of freedom.

3. Check if the **p-value** is below a predetermined significance level.

Pairwise input similarity

- The most common initialization technique if $\mathbf{A}^{(0)}$ is not given.
- **Assumption:** similar inputs should be connected.
- Input similarity can be defined in different ways. For example:
 1. Cosine similarity ($\frac{\mathbf{x}_i \cdot \mathbf{x}_j}{\|\mathbf{x}_i\| \|\mathbf{x}_j\|}$)
 2. Decreasing function of a distance \mathbf{d} (e.g., $\frac{1}{\mathbf{d}(\mathbf{x}_i, \mathbf{x}_j)}$)
 3. Kernels (e.g., the RBF kernel: $e^{-\|\mathbf{x}_i - \mathbf{x}_j\|^2}$)

😊 Easy to implement.

😊 Computationally and memory efficient.

😢 If $\mathbf{A}^{(0)}$ is not perfected afterwards, performance on the considered task may not exceed that of a structure agnostic baseline [3].

[3] Errica, “On class distributions induced by nearest neighbor graphs for node classification of tabular data” 2024.

Dirichlet Energy Minimization

- Graph signal processing perspective. [4], [5]
- Often considers symmetric and non-negative matrices. [6]
- Smoothness assumption: in amenable graph structures the graph signal varies smoothly across edges.

Define the Dirichlet Energy:

$$\mathcal{E} = \frac{1}{2} \sum_{i,j} \mathbf{A}_{ij} \|\mathbf{X}_i - \mathbf{X}_j\|^2 \equiv \frac{1}{2} \sum_{i,j} \mathbf{A}_{ij} \mathbf{Z}_{ij}$$

Minimization problem for smooth signals:

$$\mathbf{A}^{(0)} = \operatorname{argmin}_{\mathbf{A}} \left\{ \frac{1}{2} \sum_{i,j} \mathbf{A}_{ij} \mathbf{Z}_{ij} \right\}$$

Q: What is the trivial solution of this minimization problem?

[4] Dong et al., “Learning Laplacian matrix in smooth graph signal representations” 2016.

[5] Dong et al., “Learning graphs from data: A signal representation perspective” 2019.

[6] Kalofolias, “How to learn a graph from smooth signals” 2016.

Dirichlet Energy Minimization

- An additional term $f(\mathbf{A})$ imposes **prior information** and avoids converging towards the **trivial solution**.
- The complete minimization problem becomes:

$$\mathbf{A}^{(0)} = \operatorname{argmin}_{\mathbf{A}} \left\{ \frac{1}{2} \sum_{i,j} \mathbf{A}_{ij} \mathbf{Z}_{ij} + \lambda f(\mathbf{A}) \right\}$$

- 😊 The Dirichlet Energy Minimization problem provides a **theoretical framework** to different input similarity techniques. For example, if:

$$f(\mathbf{A}) = 2 \frac{\sigma^2}{\lambda} \sum_{ij} \mathbf{A}_{ij} (\log(\mathbf{A}_{ij}) - 1)$$

the solution to the minimization problem is a RBF initialization $\mathbf{A}_{ij}^{(0)} = e^{-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}}$

- 😊 Interpretable assumptions embedded in f
 😊 Rich literature present
 ☹ Less straightforward to **implement** (and **optimize**)

Rewiring techniques

- GNNs suffer from **oversmoothing** and **oversquashing** [7]
- Rewiring modifies the initial connectivity $\mathbf{A}^{(0)}$ to alleviate those problems. [8]

Oversmoothing: repeated rounds of message passing make node representations converge to similar embeddings.

Q: Connect the Dirichlet energy to oversmoothing: how does it change adding more GNN layers?

[7] Rusch *et al.*, “A survey on oversmoothing in graph neural networks” 2023.

[8] Attali *et al.*, “Rewiring Techniques to Mitigate Oversquashing and Oversmoothing in GNNs: A Survey” 2024.

Rewiring techniques

Oversquashing: exponential loss of information increases with the number of GNN layers employed.

Notation:

- $h_i^{(\ell)}$: representation of node i at layer ℓ .
- \hat{A} : normalized augmented adjacency matrix.

Given two nodes i and j at distance r , it has been shown [9]:

$$\left| \frac{\partial h_i^{(r)}}{\partial x_j} \right| \leq (K)^r (\hat{A}^r)_{ij} \quad \text{with } K \text{ being a GNN-specific constant}$$


- 😊 Changing the graph structure can alleviate both.
- [9] proposes to iteratively add and remove edges via the [Stochastic Discrete Ricci Flow](#) algorithm.
 - Some rewiring techniques completely ignore the original structure [10].

[9] Topping et al., “Understanding over-squashing and bottlenecks on graphs via curvature” 2021.

[10] Attali et al., “Delaunay Graph: Addressing Over-Squashing and Over-Smoothing Using Delaunay Triangulation” 2024.

Edge Scorer

General GSL Framework



Edge Scorer

- An **edge scorer** is a parametric function $\xi_{\theta}(\mathbf{X}, \mathbf{A})$ that returns relational structures Φ , often modeled as pairwise scores between inputs.
- Edge Scorer's parameters θ can be **trained** on the considered downstream task.

An edge scorer should:

- align, whereas possible, with **physical model**: Are scores input-dependent? Should complex relationships be considered?
- be designed having in mind **constraints** set by the problem. How many nodes are present? How much data is available?

Edge Scorer's parameters can often be initialized using extracted adjacency matrices.


Lookup table

Assume a **fixed** and **input-independent** graph structure $\longrightarrow \xi_\theta(\mathbf{X}, \mathbf{A}) = \xi_\theta$.

$N \times N$ table

The function ξ_θ is a table of parameters:


$$\xi_\theta = \Phi \in \mathbb{R}^{N \times N}$$



Embedding factorization

Parameters contained in node embeddings:



$$\xi_\theta = \Phi = Z_s Z_t^T \text{ with } Z_s, Z_t \in \mathbb{R}^{N \times d}$$



😊 Finer control

😊 More parameter efficient

Lookup table

 $N \times N$ table**Embedding factorization**

- 😊 Common choice in the literature [11]–[14]
- 😊 Easy to implement and learn
- 😢 May oversimplify the problem

[11] Franceschi *et al.*, “Learning discrete structures for graph neural networks” 2019.

[12] Wu *et al.*, “Graph wavenet for deep spatial-temporal graph modeling” 2019.

[13] Cini *et al.*, “Sparse Graph Learning from Spatiotemporal Time Series” 2023.

[14] Manenti *et al.*, “Learning Latent Graph Structures and their Uncertainty” 2024.

Input dependent

The Edge Scorer $\xi_{\theta}(\mathbf{X}, \mathbf{A})$ is a function, enabling different **inductive biases** [2], [15], [16]:

- Some methods simply use a **MLP**
- Some others employ a **Graph Neural Networks**
- Others use simple **attention-based** architectures

 Iterative score updates and GNN processing blur the distinction between the Edge Scorer and GNN. In those scenarios, a clear decomposition may not be possible.

As a general rule: **keep things simple!**


[2] Fatemi *et al.*, “Ugsl: A unified framework for benchmarking graph structure learning” 2023.

[15] Wang *et al.*, “Dynamic graph cnn for learning on point clouds” 2019.

[16] Kazi *et al.*, “Differentiable graph module (dgm) for graph convolutional networks” 2022.

Post-processing techniques & Loss functions

General GSL Framework



Post-processing techniques

The score matrix Φ is transformed into an adjacency matrix \tilde{A} to enforce desired properties.

Common objectives include:

- Training facilitation: row normalization, value clamping, etc.
- Enforcement of structures: symmetrization, minimum spanning tree construction, etc.
- Sparsification: top-k selection, Bernoulli sampling, thresholding, etc.


Specific application requirements often necessitate post-processing techniques.



Post-processing can introduce unwished consequences.

Let's focus on sparsification techniques, as it is a desirable property.

Sparse matrices



- A **sparse** matrix is a matrix in which the majority of elements are **zero**.
- **Sparsity** of a matrix = **percentage** of zero elements.

Q: Why do you think sparse matrices are desirable?

- Most common sparse representation of adjacency matrices in GDL is the COO (coordinate) format: two tensors, one for non-zero **indices** location and the other for corresponding **values**:
e.g., **indices** = [[0, 3, 5], [2, 1, 5]] **values** = [0.9, 0.9, 0.5]
- Other possibilities: CSR, CSC, BSR, BSC, ... formats

Sparse Matrices


Q: What is the computational complexity of a **dense** GCN layer $\mathbf{X}' = \mathbf{A}\mathbf{X}\mathbf{W}$?

Q: What is the computational complexity of the same GCN layer with **sparse** matrix multiplications?

Two post-processing techniques that enforce sparsity:

Thresholding


Keep edges if score > threshold



	0.1	0.9		0.1		
		0.1	0.1	0.1	0.1	
		0.1	0.1	0.1	0.1	
		0.9	0.1	0.1	0.1	0.1
	0.1		0.1	0.1	0.1	0.1
		0.1	0.1			0.5


Bernoulli sampling

Treat scores as logits to sample from



Thresholding

- Thresholding involves selecting a **threshold** hyperparameter τ and zeroing entries for which $\Phi_{ij} < \tau$.




Q: Why is the gradient biased?

- Other sparsification methods, such as top-k or top-p selection, exhibit similar advantages and disadvantages.

Bernoulli sampling

- Sample each edge with probability Φ_{ij} (or sigmoid(Φ_{ij})).



- Offers an inherently probabilistic framework.
- Gradient propagation in stochastic operations - e.g. VAEs - is challenging. In VAEs problem was solved with the reparameterization trick.
- Issues arise as gradients are computed with respect to Φ :

$$\nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X})]$$

Reparameterization trick

- Direct **sampling** from a distribution (e.g., Gaussian) introduces a **non-differentiable** operation, blocking gradient flow.
 - Reparameterization trick solves this problem **separating** the **stochastic nature** from the trainable **parameters**
1. Express the **sampled** variable \hat{A} as a **deterministic function** of trainable **parameters** Φ and a **random variable** ϵ .
 2. Example (**Gaussian**): $\hat{A} = \mu(\Phi) + \sigma(\Phi) \odot \epsilon$, where $\epsilon \sim \mathcal{N}(0, \mathbf{I})$.
 $\mu(\Phi)$ represents the mean tensor, parameterized by Φ .
 $\sigma(\Phi)$ represents the standard deviation tensor, parameterized by Φ .
 \odot is the element wise multiplication.
- ⌚ Being Bernoulli random variables discrete, the reparameterization is not applicable.

Bernoulli Sampling

- Issue arises as gradients are calculated with respect to Φ , the parameter vector defining the distribution:

$$\nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X})]$$

- Different possible gradient estimators for Bernoulli Random Variables [17]:
 1. **Straight-Through** gradient estimator (treat discrete sample as identity in backward pass) [18]
 2. **Gumbel-Softmax trick** (continuous relaxation of Bernoulli) [19]
- :(Both methods need **dense** computation or **biased** gradient estimation.
 3. **REINFORCE** and/or **Score-Function** gradient estimator. [20], [21].

[17] Mohamed *et al.*, “Monte carlo gradient estimation in machine learning” 2020.

[18] Bengio *et al.*, “Estimating or propagating gradients through stochastic neurons for conditional computation” 2013.

[19] Jang *et al.*, “Categorical Reparametrization with Gumble-Softmax” 2017.

[20] Williams, “Simple statistical gradient-following algorithms for connectionist reinforcement learning” 1992.

[21] Sutton *et al.*, “Policy gradient methods for reinforcement learning with function approximation” 1999.

Bernoulli Sampling - REINFORCE

The score function gradient estimator directly approximates the gradient of an expectation by leveraging the **log-likelihood trick** to enable gradient computation through discrete random variables.

$$\begin{aligned}\nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X})] &= \nabla_{\Phi} \int \mathcal{L}(A, \mathbf{X}) P_{\Phi}(A) dA \\ &= \int \mathcal{L}(A, \mathbf{X}) \nabla_{\Phi} P_{\Phi}(A) dA \\ &= \int \mathcal{L}(A, \mathbf{X}) P_{\Phi}(A) \nabla_{\Phi} \log P_{\Phi}(A) dA \\ &= \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X}) \nabla_{\Phi} \log P_{\Phi}(A)] \\ &\approx \frac{1}{N} \sum_{i=1}^N \mathcal{L}(A_i, \mathbf{X}) \nabla_{\Phi} \log P_{\Phi}(A_i)\end{aligned}$$

- 😊 Sparse computations and unbiased gradient estimates.
- 😢 High variance (slow or no convergence). It can be mitigated using control variates.

Bernoulli Sampling - REINFORCE

- Control variates are used to **reduce the variance** of the gradient estimate.
- **Idea:** subtract a function with known expectation from the noisy estimate.

How it works:

1. Let $\nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X})]$ be the gradient to estimate.
2. Find a control variate $c(A, \mathbf{X})$ with known expectation $\mathbb{E}_{A \sim P_{\Phi}} [c(A, \mathbf{X})]$.
3. Modify the function:

$$\nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X})] \approx \nabla_{\Phi} \mathbb{E}_{A \sim P_{\Phi}} [\mathcal{L}(A, \mathbf{X}) - \beta (c(A, \mathbf{X}) - \mathbb{E}_{A \sim P_{\Phi}} [c(A, \mathbf{X})])]$$

- ⚠ The control variate $c(A)$ should be **correlated** with $\mathcal{L}(A, \mathbf{X}) \nabla_{\Phi} \log P_{\Phi}(A)$.
- ⚠ The expectation $\mathbb{E}_{A \sim P_{\Phi}} [c(A)]$ must be known or easily computable.

Loss functions

Total loss typically composed of two components:

1. **(Un/Self-)Supervised Loss:** Drives learning towards meaningful graph structures for solving a specific downstream task.
2. **Regularization Loss:** Enforces desired properties and constraints on the learned graph.

(Self-)Supervised Loss	Regularization Loss
Downstream task (MAE, MSE, Cross-Entropy, ...)	Closeness to initial graph structure
Denoising loss	Large weights penalization (L1, L2)
Contrastive loss	Discourage large / low degree nodes Enforce symmetry Enforce or discourage specific graph density

Conclusions

Conclusions

- Learning relational structures offers a **powerful alternative** to rely on pre-defined or potentially flawed adjacency matrices
- We explored a range of techniques. Each offers different trade-offs in terms of **complexity**, **expressiveness**, and **gradient estimation properties**.

Some bits of advice:

- Don't underestimate pre-processing! If possible, **initialize** your scores.
- While challenging, try to **visualize** small learned graphs. Do the learned connections make sense in your domain?
- GSL papers are **noisy**! Check if the claims made are sustained in practice with **rigorous validations**.

Thank you for your attention!

Questions?

References i

- [1] Z. Zhiyao, S. Zhou, B. Mao, *et al.*, “**Opengsl: A comprehensive benchmark for graph structure learning**,” *Advances in Neural Information Processing Systems*, vol. 36, 2024.
- [2] B. Fatemi, S. Abu-El-Haija, A. Tsitsulin, *et al.*, “**Ugsl: A unified framework for benchmarking graph structure learning**,” *arXiv preprint arXiv:2308.10737*, 2023.
- [3] F. Errica, “**On class distributions induced by nearest neighbor graphs for node classification of tabular data**,” *Advances in Neural Information Processing Systems*, vol. 36, 2024.
- [4] X. Dong, D. Thanou, P. Frossard, and P. Vandergheynst, “**Learning laplacian matrix in smooth graph signal representations**,” *IEEE Transactions on Signal Processing*, vol. 64, no. 23, pp. 6160–6173, 2016.
- [5] X. Dong, D. Thanou, M. Rabbat, and P. Frossard, “**Learning graphs from data: A signal representation perspective**,” *IEEE Signal Processing Magazine*, vol. 36, no. 3, pp. 44–63, 2019.
- [6] V. Kalofolias, “**How to learn a graph from smooth signals**,” in *Artificial intelligence and statistics*, PMLR, 2016, pp. 920–929.
- [7] T. K. Rusch, M. M. Bronstein, and S. Mishra, “**A survey on oversmoothing in graph neural networks**,” *arXiv preprint arXiv:2303.10993*, 2023.

References ii

- [8] H. Attali, D. Buscaldi, and N. Pernelle, “**Rewiring techniques to mitigate oversquashing and oversmoothing in gnns: A survey,**” *arXiv preprint arXiv:2411.17429*, 2024.
- [9] J. Topping, F. Di Giovanni, B. P. Chamberlain, X. Dong, and M. M. Bronstein, “**Understanding over-squashing and bottlenecks on graphs via curvature,**” *arXiv preprint arXiv:2111.14522*, 2021.
- [10] H. Attali, D. Buscaldi, and N. Pernelle, “**Delaunay graph: Addressing over-squashing and over-smoothing using delaunay triangulation,**” in *Forty-first International Conference on Machine Learning*, 2024.
- [11] L. Franceschi, M. Niepert, M. Pontil, and X. He, “**Learning discrete structures for graph neural networks,**” in *International conference on machine learning*, PMLR, 2019, pp. 1972–1982.
- [12] Z. Wu, S. Pan, G. Long, J. Jiang, and C. Zhang, “**Graph wavenet for deep spatial-temporal graph modeling,**” in *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, 2019, pp. 1907–1913.
- [13] A. Cini, D. Zambon, and C. Alippi, “**Sparse graph learning from spatiotemporal time series,**” *Journal of Machine Learning Research*, vol. 24, pp. 1–36, 2023.
- [14] A. Manenti, D. Zambon, and C. Alippi, “**Learning latent graph structures and their uncertainty,**” *arXiv preprint arXiv:2405.19933*, 2024.

References iii

- [15] Y. Wang, Y. Sun, Z. Liu, S. E. Sarma, M. M. Bronstein, and J. M. Solomon, “**Dynamic graph cnn for learning on point clouds,**” *ACM Transactions on Graphics (tog)*, vol. 38, no. 5, pp. 1–12, 2019.
- [16] A. Kazi, L. Cosmo, S.-A. Ahmadi, N. Navab, and M. M. Bronstein, “**Differentiable graph module (dgm) for graph convolutional networks,**” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 2, pp. 1606–1617, 2022.
- [17] S. Mohamed, M. Rosca, M. Figurnov, and A. Mnih, “**Monte carlo gradient estimation in machine learning,**” *The Journal of Machine Learning Research*, vol. 21, no. 1, pp. 5183–5244, 2020.
- [18] Y. Bengio, N. Léonard, and A. Courville, “**Estimating or propagating gradients through stochastic neurons for conditional computation,**” *arXiv preprint arXiv:1308.3432*, 2013.
- [19] E. Jang, S. Gu, and B. Poole, “**Categorical reparametrization with gumble-softmax,**” in *International Conference on Learning Representations (ICLR 2017)*, OpenReview. net, 2017.
- [20] R. J. Williams, “**Simple statistical gradient-following algorithms for connectionist reinforcement learning,**” *Machine learning*, vol. 8, pp. 229–256, 1992.
- [21] R. S. Sutton, D. McAllester, S. Singh, and Y. Mansour, “**Policy gradient methods for reinforcement learning with function approximation,**” *Advances in neural information processing systems*, vol. 12, 1999.